

06_MNIST-CNN_v1.0-PauloBraga

July 14, 2020

Convolutional Neural Networks

```
[1]: '''
      Paulo Simplicio Braga
      11.07.2020
      '''
import warnings
warnings.filterwarnings("ignore")
import tensorflow as tf
import numpy as np
import keras
from keras.models import Sequential
from keras.layers import Conv2D, Dense, Dropout, Flatten, MaxPooling2D
print('keras using %s backend'%keras.backend.backend())
import matplotlib.pyplot as graph
import pandas as pd
%matplotlib inline
graph.rcParams['figure.figsize'] = (15,5)
graph.rcParams["font.family"] = 'DejaVu Sans'
graph.rcParams["font.size"] = '12'
graph.rcParams['image.cmap'] = 'rainbow'
```

keras using tensorflow backend

```
[2]: train_df = pd.read_csv('Data/train.csv')
      # test_df = pd.read_csv('Data/test.csv')
```

```
[3]: train_df.shape
```

```
[3]: (42000, 785)
```

```
[5]: # Linhas_Total = 42000; Linhas_60pct=25200; Linhas_20pct=8400;
      # Linhas_20pct=8400
      # train: 0~25199
      # valid: 25200~33599
      # test: 33600~41999
      # Dividindo o set de treinamento em 'train', 'validation' e 'test'
train_X = train_df.loc[:25199, train_df.columns != 'label']
```

```

train_Y = train_df.loc[:,25199, train_df.columns == 'label']

valid_X = train_df.loc[25200:33599, train_df.columns != 'label']
valid_Y = train_df.loc[25200:33599, train_df.columns == 'label']

test_X = train_df.loc[33600:41999, train_df.columns != 'label']
test_Y = train_df.loc[33600:41999, train_df.columns == 'label']

```

```

[6]: print(train_X.shape)
      print(train_Y.shape)
      print(valid_X.shape)
      print(valid_Y.shape)
      print(test_X.shape)
      print(test_Y.shape)

```

```

(25200, 784)
(25200, 1)
(8400, 784)
(8400, 1)
(8400, 784)
(8400, 1)

```

```

[7]: # Redimensiona para 28x28
      train_X = train_X.values.reshape(-1, 28, 28, 1)
      valid_X = valid_X.values.reshape(-1, 28, 28, 1)
      test_X = test_X.values.reshape(-1, 28, 28, 1)

```

```


[8]: graph.imshow(train_X[3, :, :, 0]\
                  , cmap='gray', interpolation='nearest')

```

```

[8]: <matplotlib.image.AxesImage at 0x7fc73632b390>

```



```
[9]: # Normalização
train_X = train_X/255
valid_X = valid_X/255
test_X = test_X/255
```

```
[10]: # One-hot encoding - Utilizando a função 'to_categorical', para
# converter para binário e assim a rede neural entende como
# categoria
train_Y = keras.utils.to_categorical(train_Y, 10)
valid_Y = keras.utils.to_categorical(valid_Y, 10)
test_Y = keras.utils.to_categorical(test_Y, 10)

print(train_Y[3])
```

```
[0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
```

```
[11]: # Sets a randomisation seed for replicatability.
np.random.seed(6)

# Declara o objeto
model = Sequential()
```

```
[12]: # Pré-processamento feito pela rede neural
model.add(Conv2D(28, kernel_size = (3, 3), activation = 'relu', input_shape = (28, 28, 1)))
model.add(Conv2D(56, (3, 3), activation = 'relu'))
```

```
[13]: # Pooling - Reduz a amostragem e a complexidade dos dados
model.add(MaxPooling2D(pool_size = (2, 2)))

# Dropout - Remove nodes de maneira aleatória, para evitar
# overfitting
model.add(Dropout(0.125))

# flatten - Transforma os dados em vetor
model.add(Flatten())
```

```
[14]: # Dense layers - Fazem a classificação
model.add(Dense(128, activation='relu'))

# Dropout novamente
model.add(Dropout(0.25))

# A função de ativação softmax, no 'output', faz a classificação
# de cada categoria
model.add(Dense(10, activation=tf.nn.softmax))

# Compila o modelo
model.compile(loss='categorical_crossentropy', optimizer='Adamax',\
              metrics=['accuracy'])
```

```
[15]: # Treina o modelo com os dados de treinamento e validação
training_stats = model.fit(train_X, train_Y, batch_size = 128,\
                           epochs = 10, verbose = 1, validation_data = (valid_X, valid_Y))
# Por fim, é feita a avaliação do modelo
evaluation = model.evaluate(test_X, test_Y, verbose=0)
print('Test Set Evaluation: loss = %0.6f, accuracy = %0.2f'\
      %(evaluation[0], 100 * evaluation[1]))
```

Epoch 1/10

197/197 [=====] - 33s 170ms/step - loss: 0.4332 - accuracy: 0.8672 - val_loss: 0.1706 - val_accuracy: 0.9485

Epoch 2/10

197/197 [=====] - 34s 174ms/step - loss: 0.1503 - accuracy: 0.9558 - val_loss: 0.1003 - val_accuracy: 0.9695

Epoch 3/10

197/197 [=====] - 35s 179ms/step - loss: 0.1040 - accuracy: 0.9685 - val_loss: 0.0781 - val_accuracy: 0.9769

Epoch 4/10

```

197/197 [=====] - 36s 185ms/step - loss: 0.0787 -
accuracy: 0.9766 - val_loss: 0.0648 - val_accuracy: 0.9812
Epoch 5/10
197/197 [=====] - 36s 184ms/step - loss: 0.0661 -
accuracy: 0.9803 - val_loss: 0.0608 - val_accuracy: 0.9812
Epoch 6/10
197/197 [=====] - 37s 188ms/step - loss: 0.0518 -
accuracy: 0.9838 - val_loss: 0.0576 - val_accuracy: 0.9820
Epoch 7/10
197/197 [=====] - 37s 189ms/step - loss: 0.0450 -
accuracy: 0.9860 - val_loss: 0.0533 - val_accuracy: 0.9844
Epoch 8/10
197/197 [=====] - 37s 186ms/step - loss: 0.0390 -
accuracy: 0.9884 - val_loss: 0.0521 - val_accuracy: 0.9840
Epoch 9/10
197/197 [=====] - 37s 187ms/step - loss: 0.0338 -
accuracy: 0.9895 - val_loss: 0.0493 - val_accuracy: 0.9868
Epoch 10/10
197/197 [=====] - 38s 192ms/step - loss: 0.0293 -
accuracy: 0.9917 - val_loss: 0.0487 - val_accuracy: 0.9858
Test Set Evaluation: loss = 0.041454, accuracy = 98.74

```


```

[16]: # Aqui são apresentados 5 exemplos para demonstrar a aplicação e
      # o resultado satisfatório do modelo
      for i in range(5):
          sample = test_X[i].reshape(28, 28)


          graph.imshow(sample, cmap = 'gray', interpolation = 'nearest')
          graph.show()

          prediction = model.predict(sample.reshape(1, 28, 28, 1))
          print('prediction: %i' %(np.argmax(prediction)))


```




prediction: 0




prediction: 7



prediction: 7



prediction: 2



prediction: 2